CSEE4824 Course Project

Danqing Hua $< dh2604@columbia.edu >$
Jiacheng Yang $< jy2522@columbia.edu >$

December 2012

1 Introduction

In this report, we describe our work towards the course project in CSEE4824 Computer Architecture 2012 Fall. In this project, we are asked to optimize for a specific matrix operation workload and to design a computer architecture for this workload.

We focus our optimization on standard matrix multiplication algorithm. We aim to optimize our code in an architecture-sensitive way, such that we fully exploit the ILP and TLP in hardware.

We start our architecture configuration from a single-core design. After that, we find another multi-core design and compare these two options.

Our report is structured as follows. In Section 2, we introduce our software optimization techniques. Section 3 explains how we explore the design space in a systematic way. We show our experimental results on both software optimization and architecture tuning in Section 4. We conclude our report in Section 5.

2 Software Optimization

Among the four matrix operations, matrix multiplication dominate the computation cost in this workload. Thus, we follow the Amdahl's law to optimize this operation first.

Standard matrix multiplication has $O(n^3)$ complexity. Strassen algorithm [3] is a well-known algorithm which has a lower complexity of $O(n^{2.81})$. However, Strassen algorithm is not applicable to our problem. First, the test case is not big enough to cover the constant factor in Strassen algorithm. The biggest test case is of size 201, $201^3 = 2.73 * 201^{2.81}$. In order to make Strassen algorithm faster, the constant cannot be more than 2.73. The additional work to apply the Strassen algorithm, e.g. padding the matrix with 0, malloc temporary matrix and recursive function call is overwhelming. Second, Strassen is not easy to parallelize. Thread spawning are required in a child thread, which is not supported by $SESC$.

So we focus on optimizing standard matrix multiplication algorithm, but in an architecture-sensitive way.

2.1 Loop Ordering

First of all, for matrix multiplication, there are mutliple ways to order the loops. We choose the i,j,k ordering, because we can use a register to store the temporary result $r(i, j)$. Other ordering may enable us to reference elements in input matrix A or B by registers. However, in general memory write is more costly than memory read, so it's better for us to buffer the write into a register. Within this ordering, we also benefit from spatial locality in both input matrices after we transpose matrix B.

2.2 Transpose Matrix B

In i,j,k ordering, the data access to matrix B in the inner loop is not sequential. It reads B with a stride length of n. If no optimization is made, it will introduce more cache misses, even TLB misses. We transpose the matrix B in advance, so that the access to matrix B is sequential.

1

2.3 Loop Unrolling and Register Reusing

Loop unrolling has at least the following two advantages: (1) reduce branches (2) reuse registers to avoid memory access. By loading a common value used in an outer-loop in a register, we avoid fetch the data repeatedly from L1 or L2 caches [5, 1]. Combining loop unrolling and register reusing gives us a speedup of 1.5.

In our implementation, we have two layers of loop unrollings. The first is to unroll the inner-loop (indexed by k). But what is more important here is that we also unroll the outer-loop (indexed by j). Because the inner-loop is tight, we still have abundant available registers. If we use registers as another memory hierarchy [5], we can further speed up our program. Meanwhile, although the access to $a(i,k)$ in the inner-loop is not replaced by a register, since we organized the code in a compact way, the value of $a(i,k)$ is likely to be kept in a register, so a large number of memory traffic is saved.

2.4 Cache Blocking

Cache blokcing is a technique to improve cache behavior of matrix multiplication [4]. Basically, it breaks down the matrix into cache-size blocks. It uses a block as much as possible after loading it from memory. The ideal block size is the maximum block size that can hold the entire blocks in L1/L2 cache. It is given by [2],

$$2 * BlockSize^2 * wordSize = L1\ Cache\ Size \quad (1)$$

If we use a 16K L1 cache, the block size is 32. Additionally, we also like the block size to be a multiple of the cache-line size.

We implemented the cache blocking algorithm, but the experiment shows us that the saving effect of cache blocking is not significant. Instead, it performs even worse than non-blocking algorithm, even if the simulation result shows a lower L1 miss rate.

The reason for this degradation is that the miss penalty of L1 miss is too small. Even if we use a huge L2 cache with 1M cache size, the access time to this cache is only 2-3 cycles. Meanwhile, such a small miss

penalty is very likely to be hidden by instruction-level parallelism. We tried to optimize cache blocking for alleviate L2 cache miss, however, it seems that L2 miss is also the bottleneck of our performance, at least in our experiment.

In fact, cache blocking actually counteract with the advantage of register reusing. We have to write the result matrix multiples times with cache blocking. Similar observation is found in [5], it suggests to trade L1 reuse for register utilization.

We verify our idea by increasing the access time of L2 cache (instead of using the given formula). Our experiment shows that with the increased miss penalty, cache blocking becomes more important (see Section 4).

2.5 Multi-threading

We use multi-threading to achieve thread-level parallelism .We evenly partition the rows of the output matrix. Each thread works to compute the rows that are assigned to it. There is neither skewness nor race condition in such partition. Our experiment in Section 4 shows that our multi-threading implementation has a linear speed-up.

2.6 Optimizing other operations

The techniques used to optimize the other three operations are basically the same as matrix multiplication, namely loop unrolling, register reusing and multi-threading.

What's worth a mention is that the access pattern of these three operations are simple. For scaling and addition, data is read only once, so there is no temporal locality here, but we benefit from spatial locality especially when cache-line size is large.

For matrix-vector multiplication, the matrix is read once, while the vector is read multiple times. But the vector size in the large test case is only 201. Suppose one cache-line is kept for the matrix, and one cache line is left for the output vector. Our L1 cache is big enough to hold the entire vector. Thus, we do not optimize specifically for the cache behavior of matrix-vector multiplication.

2

3 Exploring the Design Space

The general principle of our exploration is that we start from simple parameters. We gradually fix more parameter and tune for an unknown parameter. Lessons from the architecture course greatly helps us reduce the search space.

3.1 Single-core

We start exploring the design space from a single-core design. Let's first assume that we've figured out the best cache configuration.

Inorder v.s Out-of-order: Our intuition is that out-of-order is much better than inoder execution because of ILP. We verify this by the following micro-benchmark on a large test case.

The result clearly demonstrates the advantage of out-of-order execution over in-order execution, even if the issue width of the out-of-order execution is only 1.

Execution	Freq	Issue width	time	power
In-order	1G	1	2527	4.6
Out-order	1G	1	594	14.3

Table 1: Comparing in-order and out-of-order

The power of out-of-order execution is much higher than in-order execution. In-order execution consumes less power because most of its time is spent on stall. From the detailed power consumption, we can tell that a large part of its power is spent on clock instead of on execution. So using in-order execution actually wastes more power. Thus, we adhere our following exploration to out-of-order execution.

Issue width: Issue width influences the instruction-level parallelism we can achieve. So we guess that the larger is the issue width, the better the performance is.

On an moderate size test case, we fix all other parameters and vary the issue width. The result in Figure 3.1 shows that as the issue width increases, the execution time decreases. But after issue width increases over 4, the execution time stops improving on all types of cores. We conclude that the instruction-level parallelism is fully exploited in such cases. Continuing increasing issue width will only increase the cost of maintenance, such as the access time to the re-order buffers.


Figure 1: Simulation time V.S issue width

Type of cores: Each type of core varies by their number of execution units and ROB size. At a first glance, large core may be the first choice, as it carries more powerful functional units. However, we also have to lower the frequency to meet the power constraint. There is no clear good or bad in choosing the type of cores. Each core combining with appropriate issue width and frequency may meet the power constraint while preserve a good performance. But some preliminary experiments let us prefer SmallCore and MidCore because they usually have a competitive performance within a reasonable power budget.

We dismiss the idea for a heterogeneous multi-core architecture in this project, because the workload is partitioned into equal size. It doesn't make sense to bias one thread or another by giving it a more powerful core.

Frequency The guideline for tuning frequency is: as high as possible within the power constraint. So we usually tune the frequency after all other parameters are fixed.

3.2 Multi-core

Our experiment in Section 4 shows that our program has a linear speed-up. So introducing more cores, presumably is better. But considering the power constraint, we focus on a dual-core architecture to illustrate what is new in a multi-core architecture.

An interesting observation is about the power efficiency of multi-core architecture. We measure the power efficiency of an architecture as Performance per Watt. Performance is measured by KIPS (kilo-instructions-per-second) from $SESC$'s result. We plot the power efficiency of architectures with different number of cores. We can see in Figure 2 that as the number of cores increases, the power efficiency also increases.


Figure 2: Power efficiency V.S number of cores

3.3 Configuring Cache

3.3.1 Cache Hierarchy

For single core architecture, a two-level cache hierarchy is a better choice. L1 cache guarantees fast access, while the larger L2 cache stores more data to avoid the penalty of memory access.

For multi-core architecture, we choose a shared-L2 hierarchy. In our workload, matrix B is shared by all threads. Using a private L2 cache will only lead to unnecessary cold start.

3.3.2 L1 Cache

Our code is very small, so we don't need a large L1 instruction cache. We fix L1 instruction cache to its minimum size 16K. L1 data cache is chosen to 16K as well, mainly because it saves power and L2 cache access is cheap.

3.3.3 L2 Cache

Unlike L1 cache, whose size has a significant influence on power. Enlarging L2 cache seems to come for free. We vary the size of L2 cache from 16K to 1M, the power remains stable, and the access time is within 3 cycles. The total size of three matrices is $201 * 201 * 3 * 8 = 946K$, so a 1M L2 cache can hold almost every data. Misses in L2 cache in this case are due to compulsory misses during matrix generation or conflict misses.

Furthermore, we notice that the size of L2 actually doesn't have much impact on our performance. Having a 16K L2 cache only leads to a 6% performance drop versus 256K L2 in our final configuration (from 100 msec to 106 msec). We assume it is because the access time to memory in the simulator is small, and the DRAM access is small due to our optimization.

3.3.4 Cache block size

Because almost all reads are sequential in this workload, we prefer a larger cache-line size. We set the block size to 64 bytes, because it guarantees a low access time for L1 cache and it won't bring in too much unnecessary data at the boundary of a row or a partition.

4 Experiment

In this section, we illustrate some experimental results to support our claims in Section 2 and Section 3.

4.1 Cache Blocking

We mentioned in Section 2 that cache blocking does not give us too much performance gain. Instead, it

hurts our performance.

First, we show the result by varying block size in a single-core configuration: 1.8 GHz MidCore, issue width = 2. The L1 cache is 16KB, and the L2 cache is 1M, block size is 64 bytes. The workload is medium. The start-up time in main the function(e.g. matrix generation) is roughly 1.9 msec.

Here, we find the following observations (1) the best block size is smaller the ideal block size we calculated before (2) in our workload, the performance downgrades with blocking. The block size should be smaller because of conflict in L1 cache.

blocksize	time	L1 miss	L2 miss	%Dram
16	6.944	0.67	12.21	0.53
24	6.682	1.01	8.54	0.80
32	6.711	1.66	5.19	1.31
40	6.585	1.89	4.70	1.48
None	6.514	1.78	5.08	1.40

Table 2: Cache blocking: Effect of block size

To see when the cache blocking algorithm works, we increase the access time of L2 cache. When the gap between L1 cache and L2 cache increases, cache blocking becomes more important. Blocking will have more effect in reducing L2 cache miss, supposing the gap between L2 and memory is large, say hundreds of cycles. However, it seems in this simulator the gap between L2 and memory is not large enough either (see above Section).

L2 access time	Ideal	Nonblocking	Speedup
3	6.682	6.514	0.97
30	6.899	7.168	1.04
50	7.082	8.036	1.13
100	7.589	10.647	1.40
200	8.624	15.888	1.84

Table 3: Cache blocking: Effect of L2 access time

4.2 Loop Unrolling

Not let's look at the effect of loop unrolling and how we choose the best number of unrollings we made based on our experiment. It is clear from the result that unrolling the j loop 2 times is the best choice. As the unrolling times reaches 3, the performance drops, because we have used up all the registers, and we will have a register thrashing here. The inner-most loop can be unrolled 4 times or 2 by changing a switch in the code.

Unrolling	time	L1 miss	L2 miss	Speedup
None	10.211	1.93	5.04	1.3
1	6.821	1.81	5.08	1.38
2	6.515	1.78	5.08	1.4
3	6.59	1.64	5.08	1.42

Table 4: Speedup of loop-unrolling

4.3 Multi-threading

We test the scalability of our multi-threading program. In this experiment, we ignore the power constraint: we just assign more cores and study the speedup.The time we report is the time without start-up time (matrix generation, etc). Our basic configuration is n MidCore with issue width 2, we vary n. Each core has 16K L1 Instruction cache and 32K L1 Data cache. The L2 shared cache is 256K. We use out-of-order execution and set the frequency to 1.8GHz. We run this test on medium.

When we scale the number of threads from 1 to 4, it has almost a linear scale-up. When the number of threads reaches 8, the speedup a sub-linear scale-up. Adding more threads will saturate common memory bus, and increases coordination cost. And we will suffer from skewness, i.e. the execution is limited by the slowest thread. So we cannot have linear scale-up with too many threads.

5

Figure 3: Scaling with more threads

Test	time	L1 miss	L2 miss	Power
Small	0.725	0.28	65.21	22.447
Medium	5.9	1.6	9.77	27.505
Large	103.247	1.64	95.82	29.848

Table 7: Multi-core final result

4.4 Final Performance

We report the final performance of our architecture and program. We present two configurations here, one for single-core architecture and the other for multi-core architecture. The configurations for both architectures are shown in Table 5. The results are reported in Table 6 and Table 7.

	Single	Multi
Core type	Mid	Small
Inorder	false	false
Issue width	2	2
Frequency	2.1G	1.6G
Cache Org	L2Single	L2Share
L1 Cache	16K	16K
L2 Cache	512K	256K
Block size	64	64

Table 5: Configurations for Single-core and Multi-core architecture

Test	time	L1 miss	L2 miss	Power
Small	0.606	0.29	56.76	25.131
Medium	5.585	1.78	5.08	27.907
Large	100.055	1.89	5.97	29.493

Table 6: Single-core final result

There is no clear winner between a single-core and multi-core design. In this application, there is abundant instruction-level parallelism: instructions and data are mostly independent, data dependency is less. With software optimization, we can exploit ILP to the extreme. Single-core architecture is also more simple and efficient. We tried to use 'L2Shared' model with one thread, and it performs much worse than the 'SingleL2' model. Although our multi-core architecture does not win over the single-core design within this power constraint, by our experiment, we show that our algorithm design is scalable with the number of cores.

5 Conclusion

In this report, we summarize our techniques to optimize matrix operations for the given workload. We use pre-processing, software loop-unrolling, register reusing to enhance the execution of single thread execution. We investigate cache-aware matrix multiplication and their impact on our system's performance. We apply multi-threading to parallelize all the operations and we experimentally show that we achieve a nearly-linear scale-up when the number of cores is less than 8.

We systematically explore the design space, finding out the best architecture for both uni-core architecture and dual-core architecture. Instead of enumerating all possibilities in the design space, our search is guided by our understanding in architecture concept, so we greatly reduce the search space.

Finally, we experimentally verify our design and quantify the speedup in our software optimization.

References

[1] `http://inst.cs.berkeley.edu/~cs61c/fa12/projects/03/`.

[2] `http://networks.cs.ucdavis.edu/~amitabha/optimizingMatrixMultiplication.pdf`.

[3] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein. *Introduction to Algorithms*. MIT Press, 2001.

[4] Monica D. Lam, Edward E. Rothberg, and Michael E. Wolf. The cache performance and optimizations of blocked algorithms. *SIGPLAN Not.*, 1991.

[5] David Parello, Olivier Temam, and Jean-Marie Verdun. On increasing architecture awareness in program optimizations to bridge the gap between peak and sustained processor performance: matrix-multiply revisited. In *Proceedings of the 2002 ACM/IEEE conference on Supercomputing*.